


RAFT Consensus

RAFT Consensus


- **Basic idea -**
 - The nodes collectively select a *leader*; others become *followers*
 - The leader is responsible for state transition log replication across the followers

Electing the Leader: Voting Request




- **term:** last calculated # known to candidate + 1
- **index:** committed transaction available to the candidate

Electing the leader: Follower Node's Decision Making



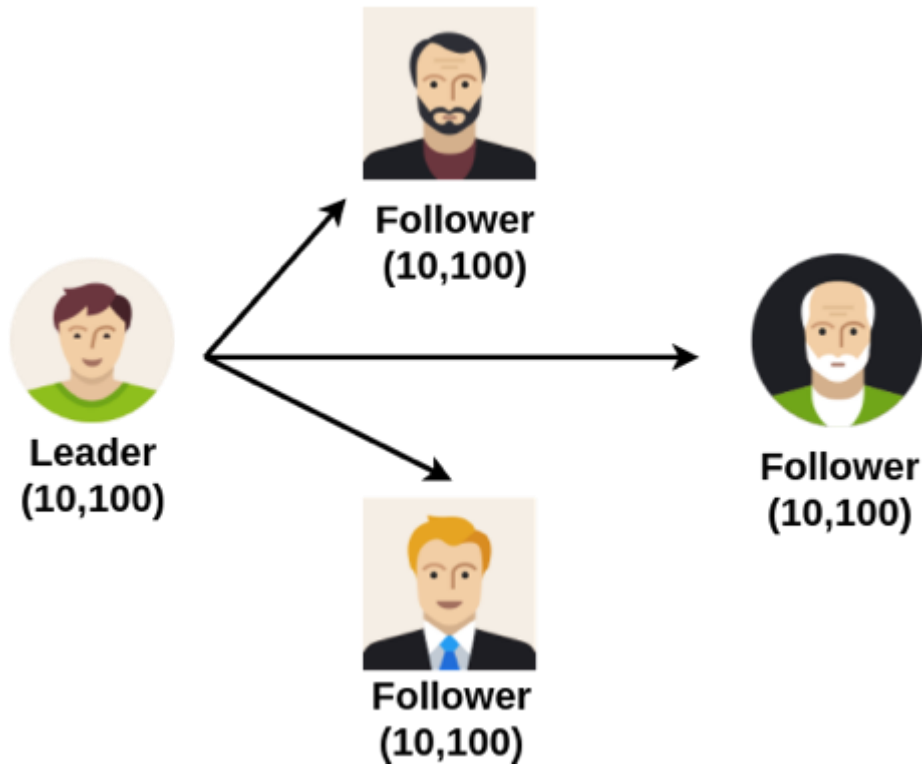
- Each node compares received term and index with corresponding current known values

Electing the leader: Majority Voting



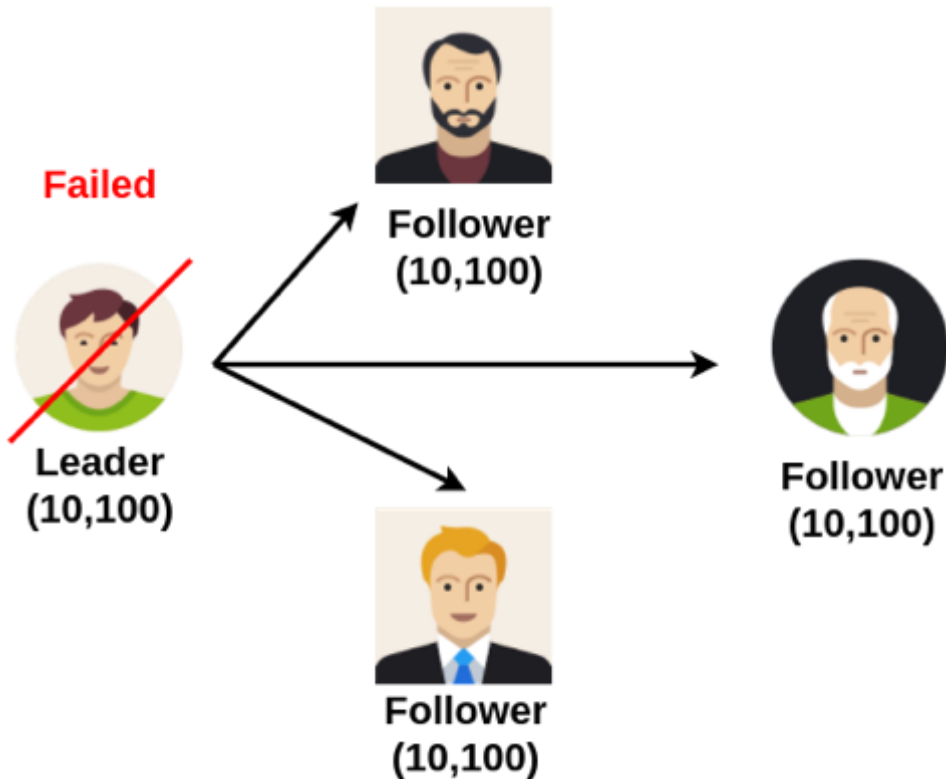
- Use of **Majority voting**
 - leader selection
 - commit the log entry

Multiple Leader Candidates: Current Leader Failure



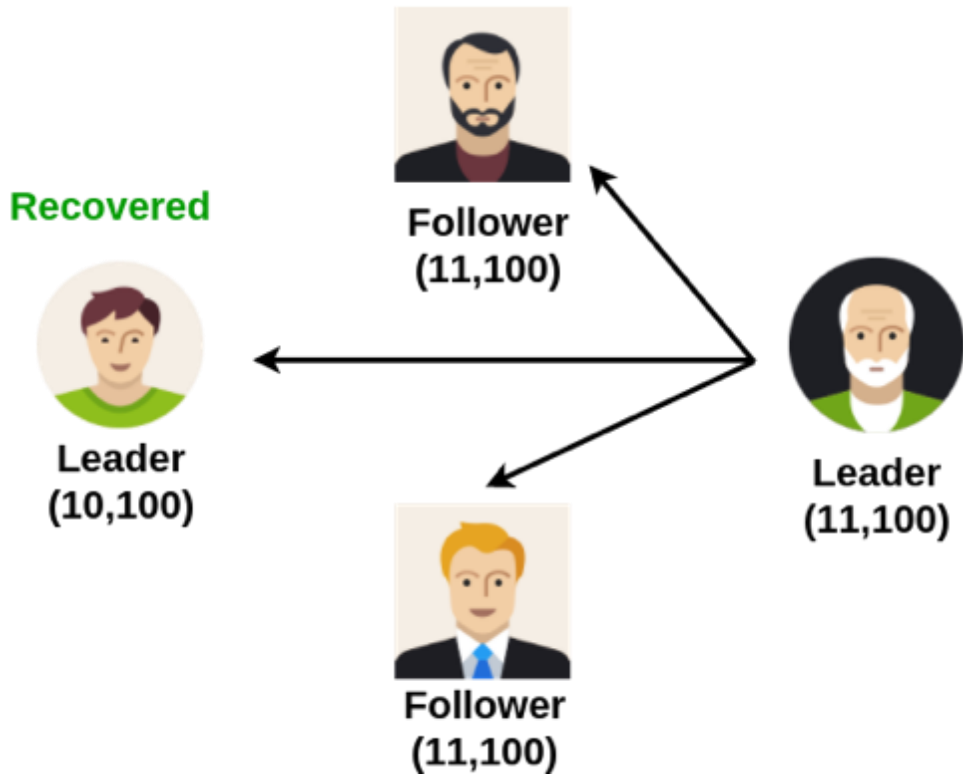
- A leader with three followers
- **term: 10**
- **commit index: 100**

Multiple Leader Candidates: Current Leader Failure



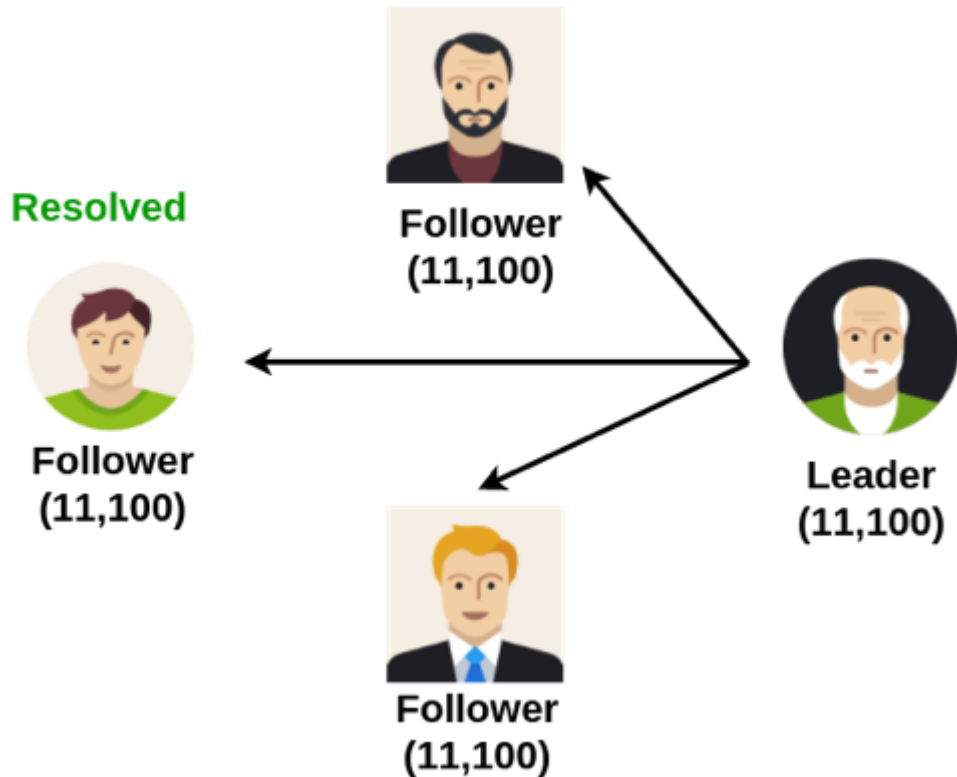
- The leader node failed

Multiple Leader Candidates: Current Leader Failure



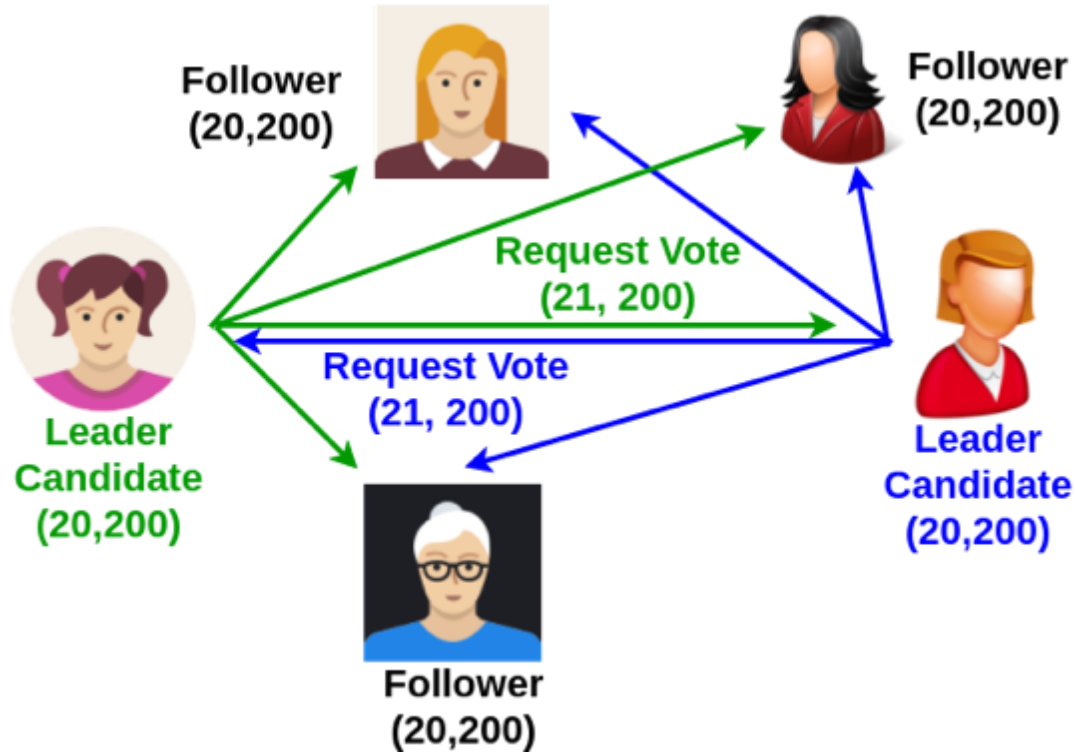
- New leader elected with term 11
- Old leader recovered

Multiple Leader Candidates: Current Leader Failure



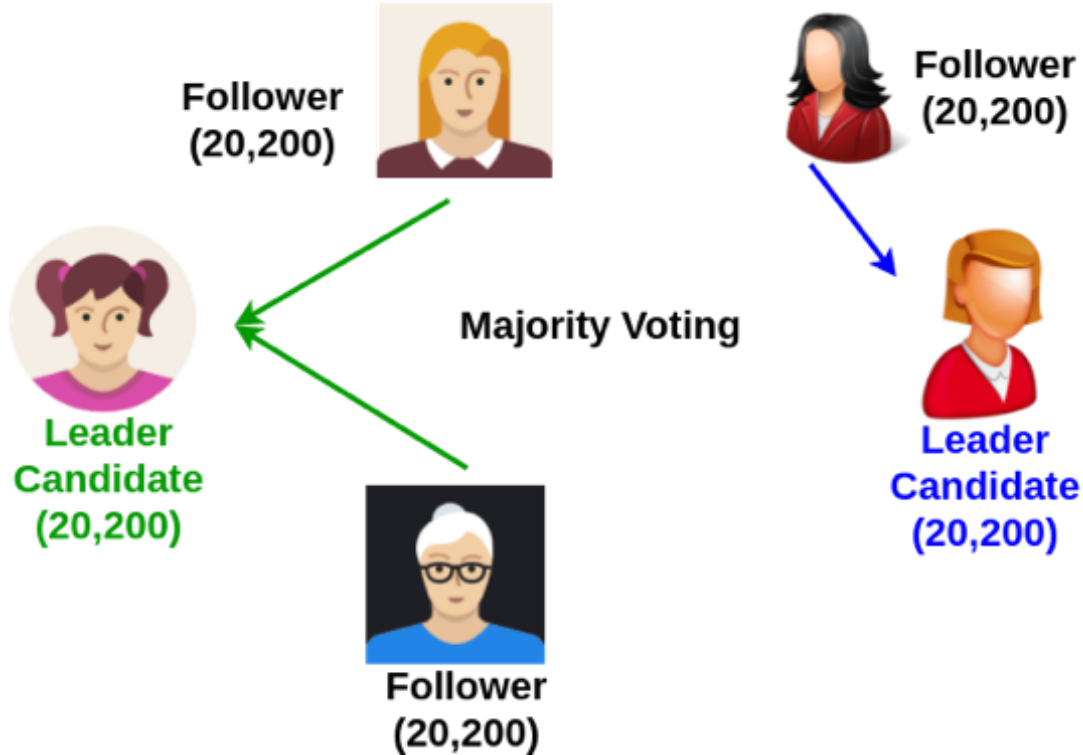
- Old leader receive heartbeat message from new leader with greater term
- Old leader drops to follower state

Multiple Leader Candidates: Simultaneous Request Vote



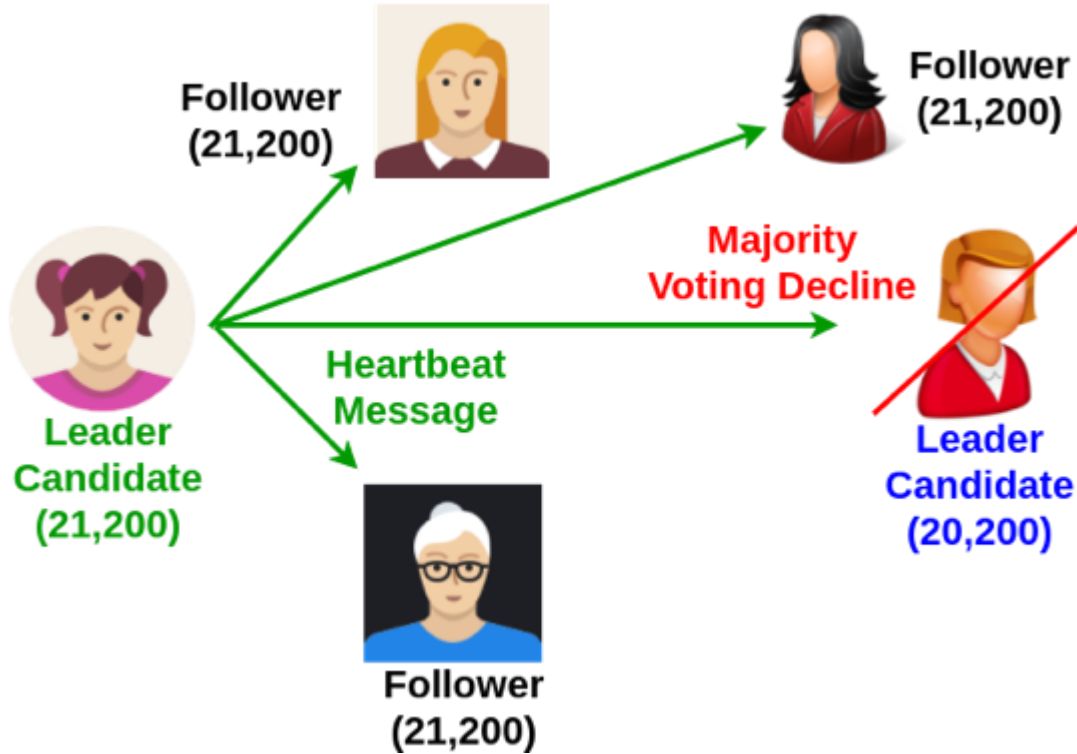
- Two nodes send Request vote message with term 21 at the same time

Multiple Leader Candidates: Simultaneous Request Vote



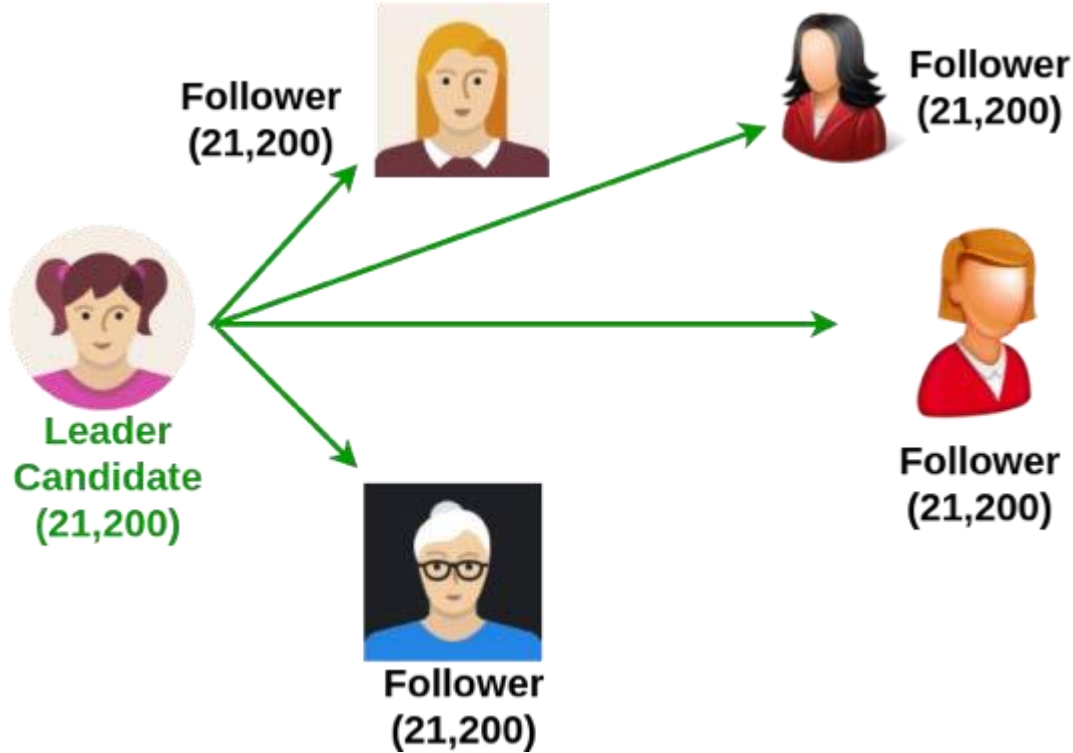
- One of them gets majority voting

Multiple Leader Candidates: Simultaneous Request Vote



- Winner sends heartbeat message

Multiple Leader Candidates: Simultaneous Request Vote



- Other leader candidate switches to follower state

Committing Entry Log



Follower
(10,100)



Leader
(10,100)

Logs: 12/3/18
12:00:00



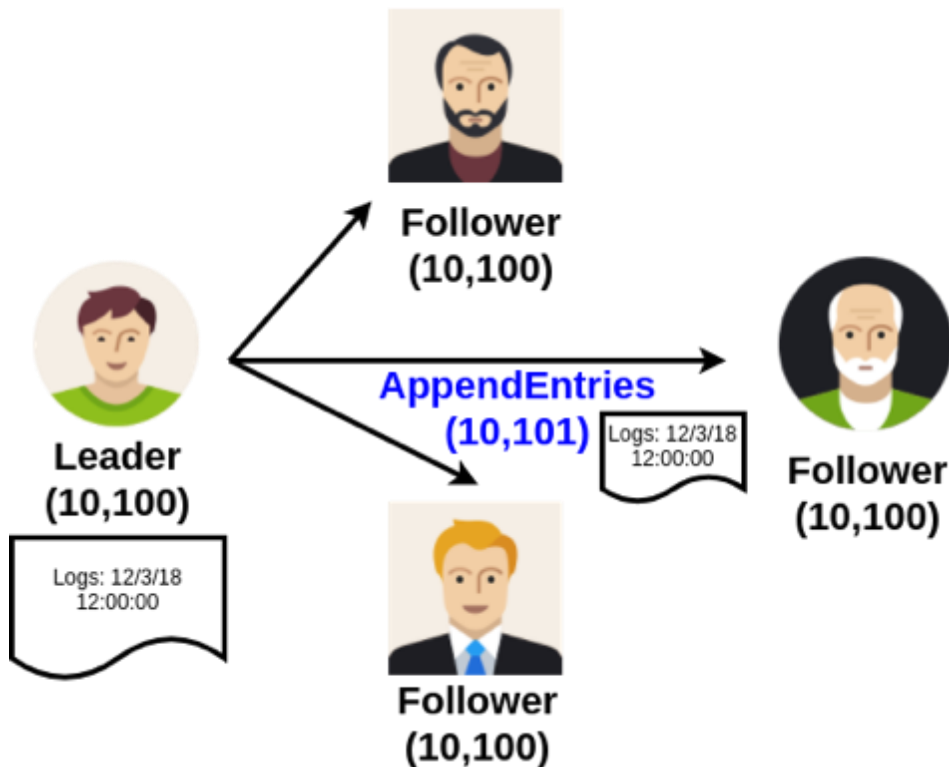
Follower
(10,100)



Follower
(10,100)

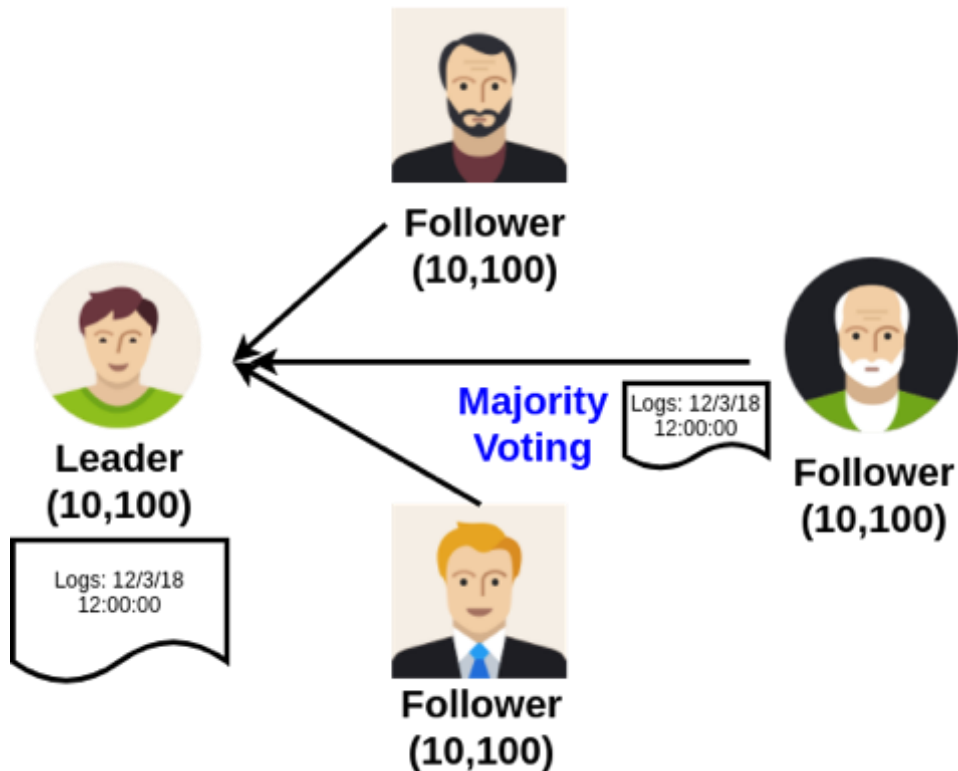
- Leader adds entry to log with term 10 and index 101

Committing Entry Log



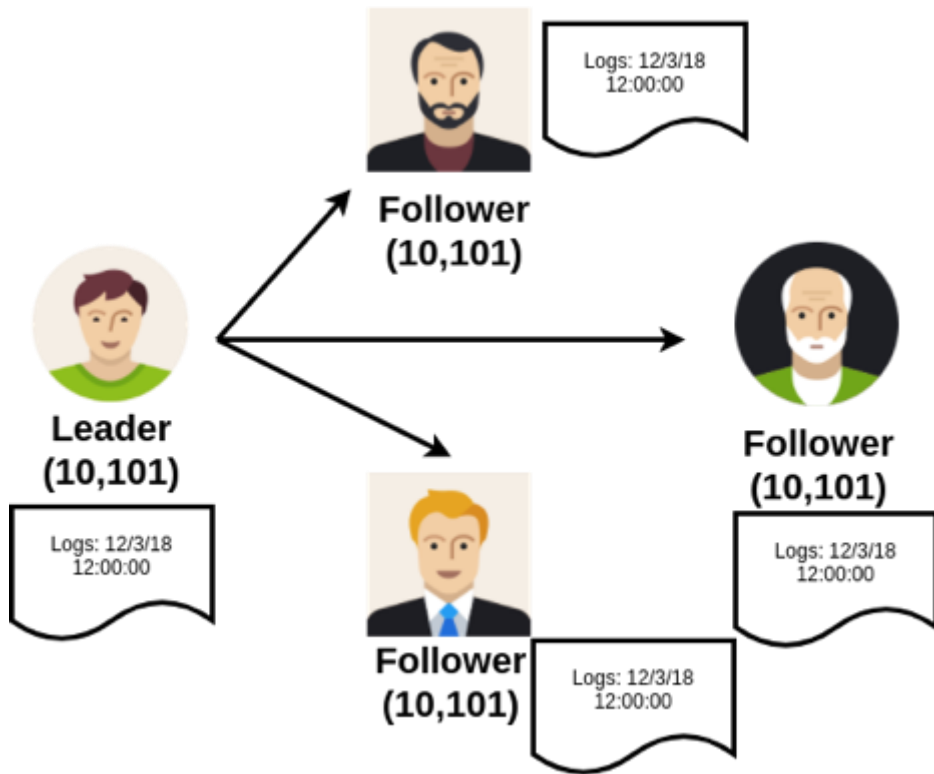
- Leader sends *AppendEntries* message to followers with index 101

Committing Entry Log



- Majority voting decides to accept or reject the entry log

Committing Entry Log



- Successfully accept entry log
 - All leader and followers update committed index to 101