



Byzantine Generals Problem

Three Byzantine Generals Problem: Lieutenant Faulty



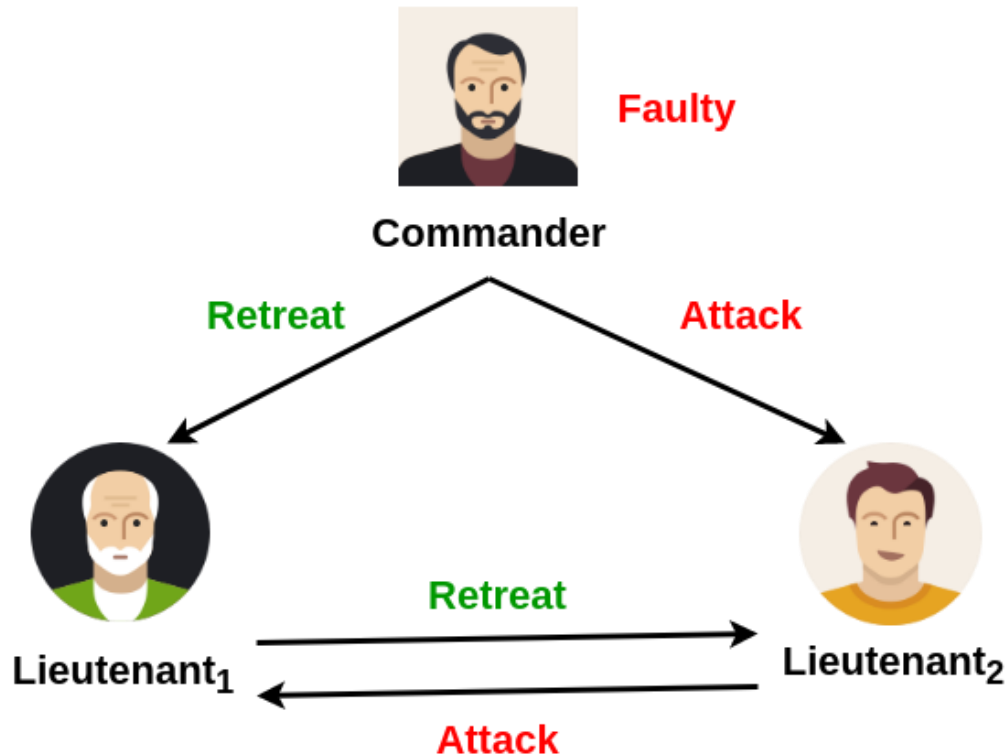
- Round1:
 - Commander correctly sends same message to Lieutenants
- Round 2:
 - Lieutenant₁ correctly echoes to Lieutenant₂
 - Lieutenant₂ **incorrectly** echoes to Lieutenant₁

Three Byzantine Generals Problem: Lieutenant Faulty




- Lieutenant₁ received **differing message**
- By integrity condition, Lieutenant₁ bound to decide on Commander message
- **What if Commander is faulty??**

Three Byzantine Generals Problem: Commander Faulty



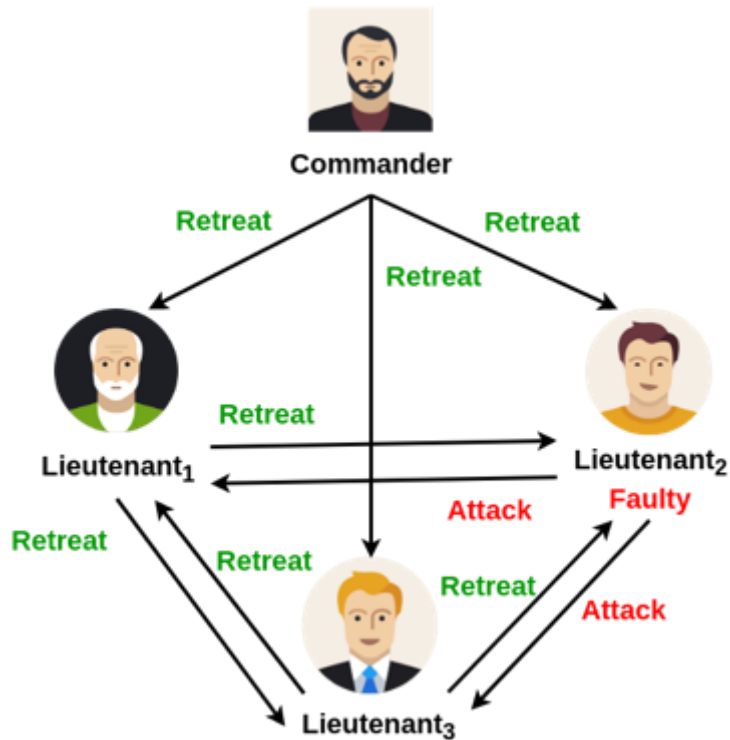
- Round 1:
 - Commander sends differing message to Lieutenants
- Round 2:
 - Lieutenant₁ correctly echoes to Lieutenant₂
 - Lieutenant₂ correctly echoes to Lieutenant₁

Three Byzantine Generals Problem: Commander Faulty



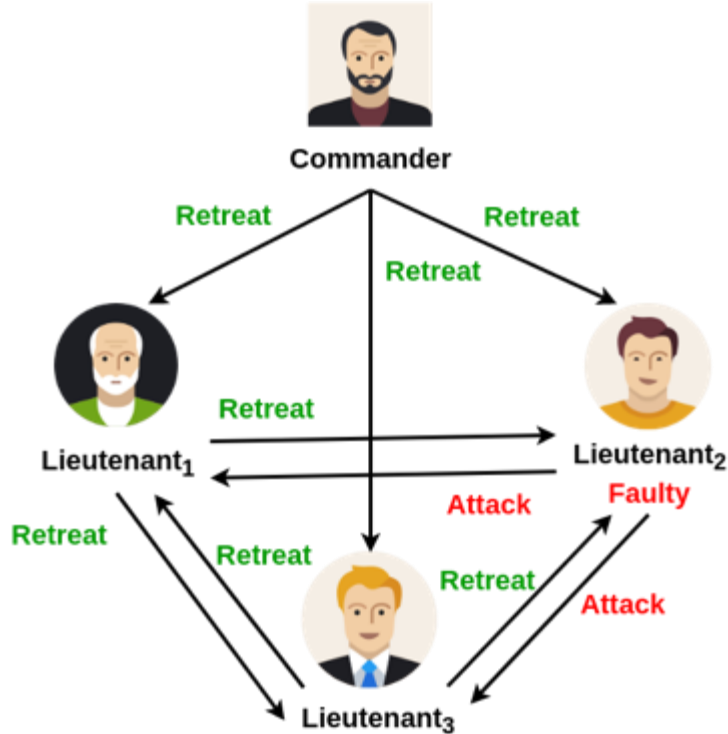
- Lieutenant₁ received **differing message**
- By integrity condition, both Lieutenants conclude with Commander's message
- This contradicts the agreement condition
- No solution possible for three generals including one faulty

Four Byzantine Generals Problem: Lieutenant Faulty




- Round 1:
 - Commander sends a message to each of the Lieutenants
- Round 2:
 - Lieutenant₁ and Lieutenant₃ correctly echo the message to others
 - Lieutenant₂ **incorrectly** echoes to others

Four Byzantine Generals Problem: Lieutenant Faulty




- Lieutenant₁ decides on $\text{majority}(\text{Retreat}, \text{Attack}, \text{Retreat}) = \text{Retreat}$
- Lieutenant₃ decides on $\text{majority}(\text{Retreat}, \text{Retreat}, \text{Attack}) = \text{Retreat}$

Four Byzantine Generals Problem: Commander Faulty




- Round 1:
 - Commander sends differing message to Lieutenants
- Round 2:
 - Lieutenant₁, Lieutenant₂ and Lieutenant₃ correctly echo the message to others

Four Byzantine Generals Problem: Commander Faulty



- Lieutenant₁ decides on $\text{majority}(\text{Retreat}, \text{Attack}, \text{Retreat}) = \text{Retreat}$
- Lieutenant₂ decides on $\text{majority}(\text{Attack}, \text{Retreat}, \text{Retreat}) = \text{Retreat}$
- Lieutenant₃ decides on $\text{majority}(\text{Retreat}, \text{Retreat}, \text{Attack}) = \text{Retreat}$


Byzantine Generals Model



- N number of process with at most f Faulty = $2f+1$
- Receiver always knows the identity of the sender
- Fully connected
- Reliable communication medium

Practical Byzantine Fault Tolerant Model

Practical Byzantine Fault Tolerant Model



- Asynchronous distributed system
 - delay, out of order message
- Byzantine failure handling
 - arbitrary node behavior
- Privacy
 - tamper-proof message, authentication


Practical Byzantine Fault Tolerant Model

- $3f + 1$ replicas are there where f is the number of faulty replicas
- The replicas move through a succession of configurations, known as *views*
- One replica in a *view* is *primary* and others are *backups*

Practical Byzantine Fault Tolerant Model


- Views are changed when a *primary* is detected as faulty
- Every view is identified by a unique integer number v
- Only the messages from the current views are accepted

Practical Byzantine Fault Tolerant Algorithm



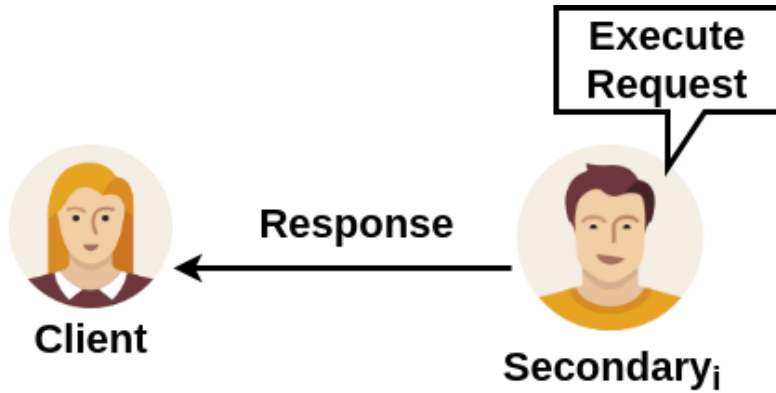
- A client sends a request to invoke a service operation to the primary

Practical Byzantine Fault Tolerant Algorithm



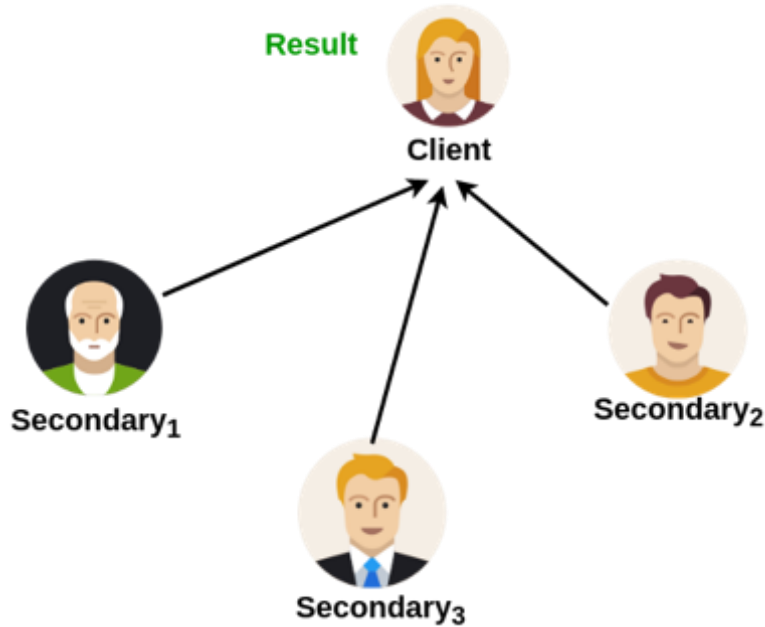
- The primary multicasts the request to the backups/Secondary nodes

Practical Byzantine Fault Tolerant Algorithm



- Backups execute the request and send a reply to the client

Practical Byzantine Fault Tolerant Algorithm




- The client waits for $f + 1$ replies from different backups with the same result
 - f is the maximum number of faulty replicas that can be tolerated

Three Phase Commit Protocol - Pre-Prepare

- **Pre-prepare:** Primary assigns a sequence number n to the request and multicast a message $\langle \langle PRE-PREPARE, v, n, d \rangle_{\sigma_p}, m \rangle$ to all the backups
 - v is the current view number
 - n is the message sequence number
 - d is the message digest
 - σ_p is the private key of primary - works as a digital signature
 - m is the message to transmit

Three Phase Protocol

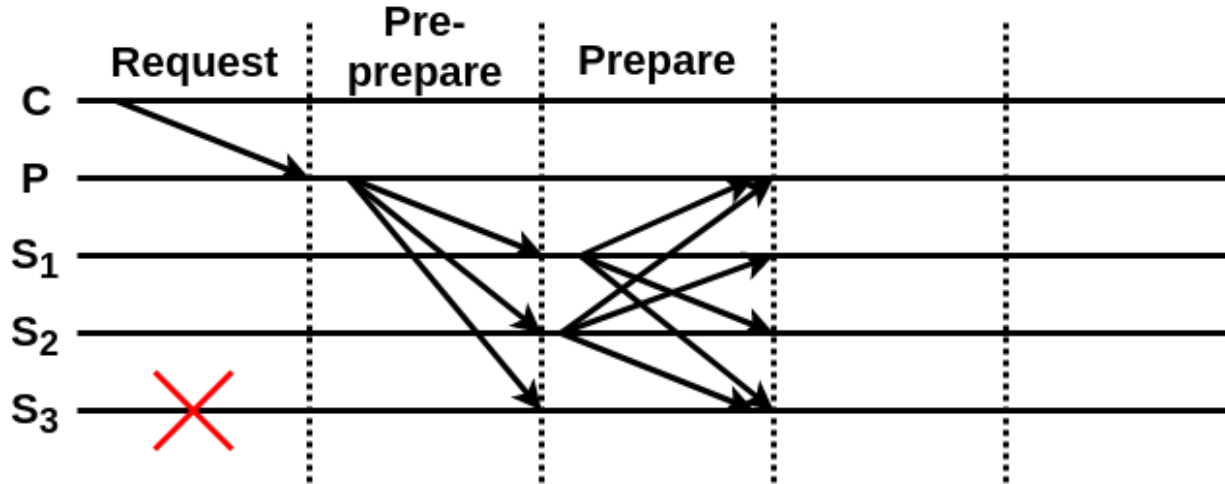


- Pre-prepare:
 - Acknowledge the request by a unique sequence number

Three Phase Commit Protocol - Pre-Prepare

- Pre-prepare messages are used as a proof that request was assigned sequence number n is the view v
- A backup accepts a pre-prepare message if
 - The signature is correct and d is the digest for m
 - The backup is in view v
 - It has not received a different PRE-PREPARE message with sequence n and view v with a different digest
 - The sequence number is within a threshold

Three Phase Protocol



- Prepare:
 - Replicas agree on the assigned sequence number

Three Phase Commit Protocol - Prepare

- If the backup accepts the PRE-PREPARE message, it enters prepare phase by multicasting a message $\langle \text{PREPARE}, v, n, d, i \rangle_{\sigma_i}$ to all other replicas
- A replica (both primary and backups) accepts prepare messages if
 - Signatures are correct
 - View number equals to the current view
 - Sequence number is within a threshold

Three Phase Commit Protocol

- Pre-prepare and prepare ensure that non-faulty replicas guarantee on a total order for the requests within a view
- Commit a message if
 - $2f$ prepares from different backups matches with the corresponding pre-prepare
 - You have total $2f + 1$ votes (one from primary that you already have!) from the non-faulty replicas

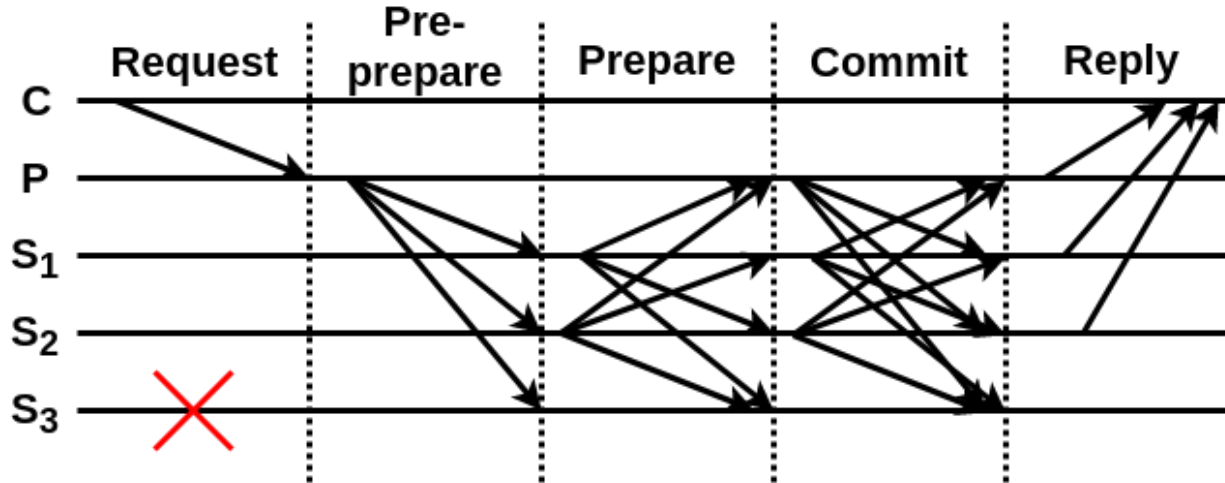
Three Phase Commit Protocol

- **Why do you require $3F+1$ replicas to ensure safety in an asynchronous system when there are F faulty nodes?**
 - If you do not receive a vote
 - The node is faulty and not forwarded a vote at all
 - The node is non-faulty, forwarded a vote, but the vote got delayed
 - Majority can be decided once $2f+1$ votes have arrived - even if f are faulty, you know $f+1$ are from correct nodes, do not care about the remaining f votes

Three Phase Commit Protocol - Commit

- Multicast $\langle \text{COMMIT}, v, n, d, i \rangle_{\sigma_i}$ message to all the replicas including primary
- Commit a message when a replica
 - Has sent a commit message itself
 - Has received $2f+1$ commits (including its own)

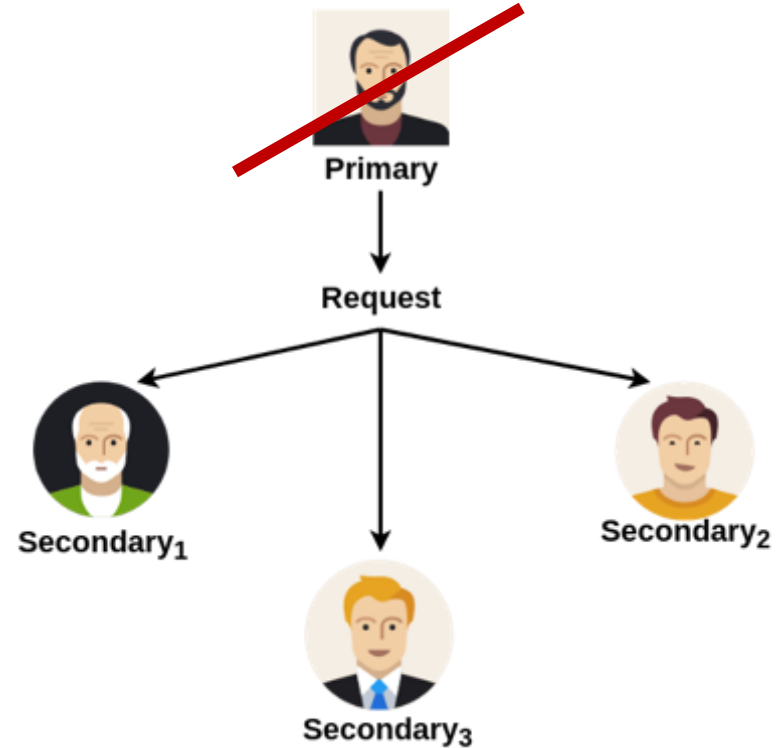
Three Phase Protocol



- Commit:
 - Establish consensus throughout the views

View Change

- What if the **primary** is **faulty**??
 - non-faulty replicas detect the fault
 - replicas together start view change operation



View Changes

- View-change protocol provides **liveness**
 - Allow the system to make progress when primary fails
- If the primary fails, backups will not receive any message (such as PRE_PREPARE or COMMIT) from the primary
- View changes are triggered by timeouts
 - Prevent backups from waiting indefinitely for requests to execute

Practical Byzantine Fault Tolerant

- **Why Practical?**
 - Ensures safety over an asynchronous network Byzantine Failure
 - Low overhead
- **Real Applications**
 - Tendermint
 - IBM's Openchain
 - ErisDB
 - Hyperledger